


Disclaimer

- These are not slides from DTU (refer to p4.org)
- You can find the original slides at:
 - https://github.com/p4lang/tutorials/blob/master/P4_tutorial.pdf
- We removed some of the slides and organized in several groups
- You are welcome to check the original slides for a wider perspective
- If you have doubts about:
 - What P4 is, why it is beneficial or the core concepts of P4: Refer to file “**1 - P4 Information.pdf**”
 - What the core elements of P4 programming are (controls, externs, tables, data types), how to program a target, how the architecture refers to a P4 program and then: Refer to file “**2 - Basics on P4 programming.pdf**”
 - What P4 runtime is, how it related to P4 overall and what you can do with it: Refer to file “**3 - P4 Runtime.pdf**”



P4 Language Tutorial



Goals

- **Learn P4 Language**
 - Traditional applications
 - Novel applications
- **Learn P4 software tools**
 - P4 Compiler
 - BMv2
 - P4Runtime
- **Learn about P4 hardware targets**
 - mini-workshop featuring solutions by Barefoot, Netronome, Netcope and NetFPGA.
- **Networking (the other kind)**
- **Have fun!**



What is Data Plane Programming?

- Why program the Data Plane?


Software Defined Networking: Logically Centralized Control

- **Main contributions**

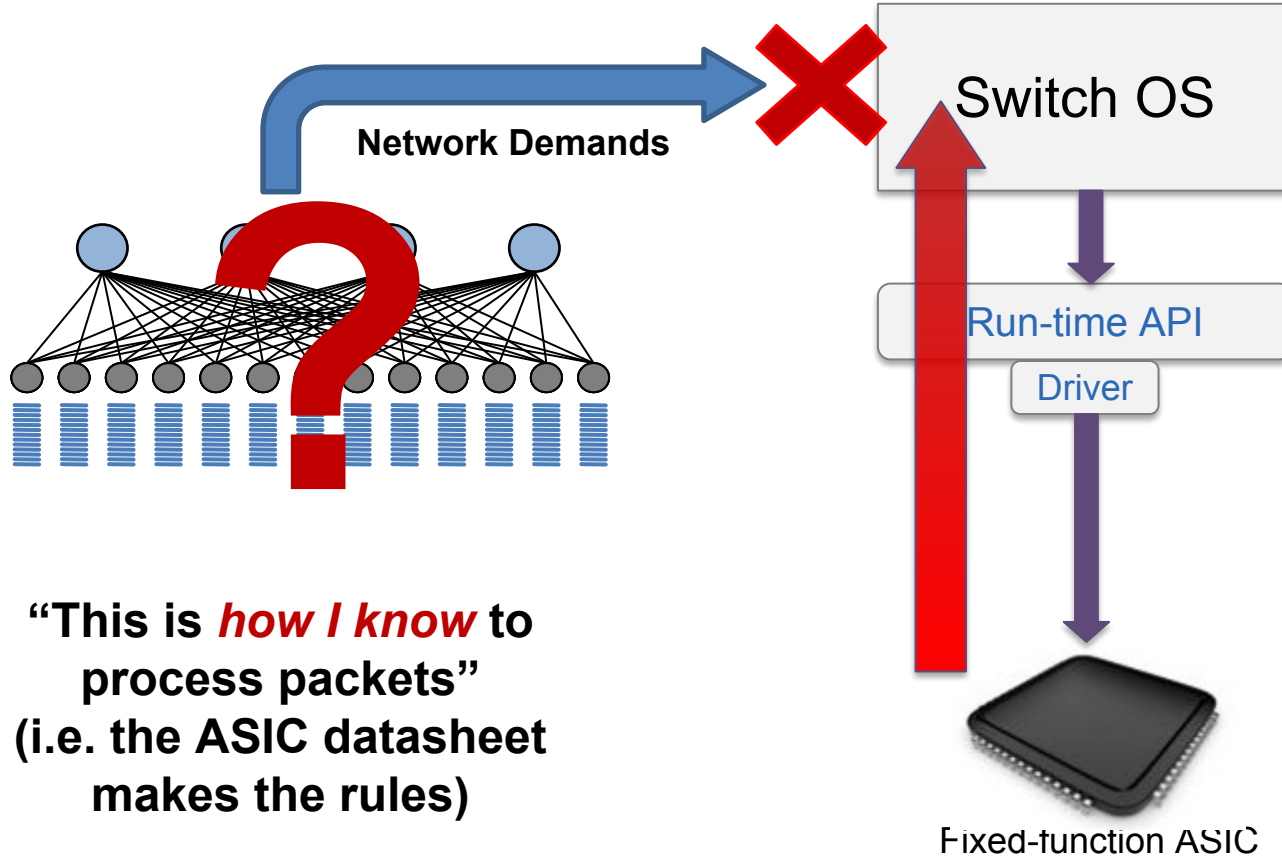
- OpenFlow = standardized *model*
 - match/action abstraction
- OpenFlow = standardized *protocol* to interact with switch
 - download flow table entries, query statistics, etc.
- *Concept* of *logically* centralized control via a single entity (“SDN controller”)
 - Simplifies control plane – e.g. compute optimal paths at one location (controller), vs. waiting for distributed routing algorithms to converge

- **Issues**

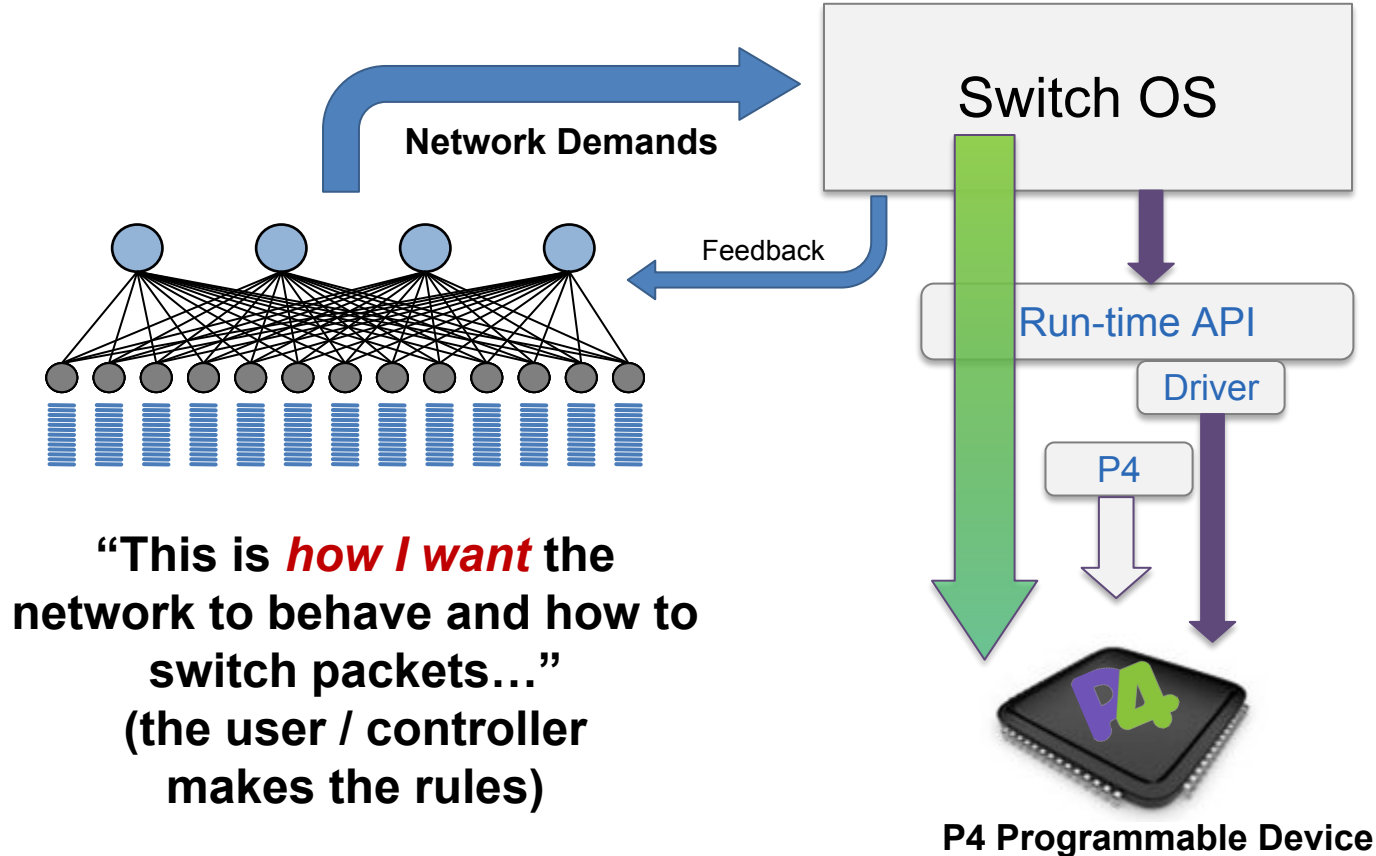
- Data-plane protocol evolution requires changes to standards (12 → 40 OpenFlow match types)
- Limited interoperability between vendors => southbound I/F differences handled at controller (OpenFlow / netconf / JSON / XML variants)



Status Quo: Bottom-up design



A Better Approach: Top-down design



Benefits of Data Plane Programmability

- **New Features** – Add new protocols
- **Reduce complexity** – Remove unused protocols
- **Efficient use of resources** – flexible use of tables
- **Greater visibility** – New diagnostic techniques, telemetry, etc.
- **SW style development** – rapid design cycle, fast innovation, fix data plane bugs in the field
- **You keep your own ideas**

Think programming rather than protocols...



Programmable Network Devices

- **PISA: Flexible Match+Action ASICs**

- Intel Flexpipe, Cisco Doppler, Cavium (Xpliant), Barefoot Tofino, ...

- **NPU**

- EZchip, Netronome, ...

- **CPU**

- Open Vswitch, eBPF, DPDK, VPP...

- **FPGA**

- Xilinx, Altera, ...

These devices let us tell them how to process packets.



What can you do with P4?

- **Layer 4 Load Balancer – SilkRoad[1]**
- **Low Latency Congestion Control – NDP[2]**
- **In-band Network Telemetry – INT[3]**
- **Fast In-Network cache for key-value stores – NetCache[4]**
- **Consensus at network speed – NetPaxos[5]**
- **Aggregation for MapReduce Applications [6]**
- **... and much more**

[1] Miao, Rui, et al. "SilkRoad: Making Stateful Layer-4 Load Balancing Fast and Cheap Using Switching ASICs." SIGCOMM, 2017.

[2] Handley, Mark, et al. "Re-architecting datacenter networks and stacks for low latency and high performance." SIGCOMM, 2017.

[4] Kim, Changhoon, et al. "In-band network telemetry via programmable dataplanes." SIGCOMM. 2015.

[3] Xin Jin et al. "NetCache: Balancing Key-Value Stores with Fast In-Network Caching." To appear at SOSP 2017

[5] Dang, Huynh Tu, et al. "NetPaxos: Consensus at network speed." SIGCOMM, 2015.

[6] Sapio, Amedeo, et al. "In-Network Computation is a Dumb Idea Whose Time Has Come." *Hot Topics in Networks*. ACM, 2017.



Brief History and Trivia

- **May 2013:** Initial idea and the name “P4”
- **July 2014:** First paper (SIGCOMM CCR)
- **Aug 2014:** First P4₁₄ Draft Specification (v0.9.8)
- **Sep 2014:** P4₁₄ Specification released (v1.0.0)
- **Jan 2015:** P4₁₄ v1.0.1
- **Mar 2015:** P4₁₄ v1.0.2
- **Nov 2016:** P4₁₄ v1.0.3
- **May 2017:** P4₁₄ v1.0.4

- **Apr 2016:** P4₁₆ – first commits
- **Dec 2016:** First P4₁₆ Draft Specification
- **May 2017:** P4₁₆ Specification released



Why P4₁₆?

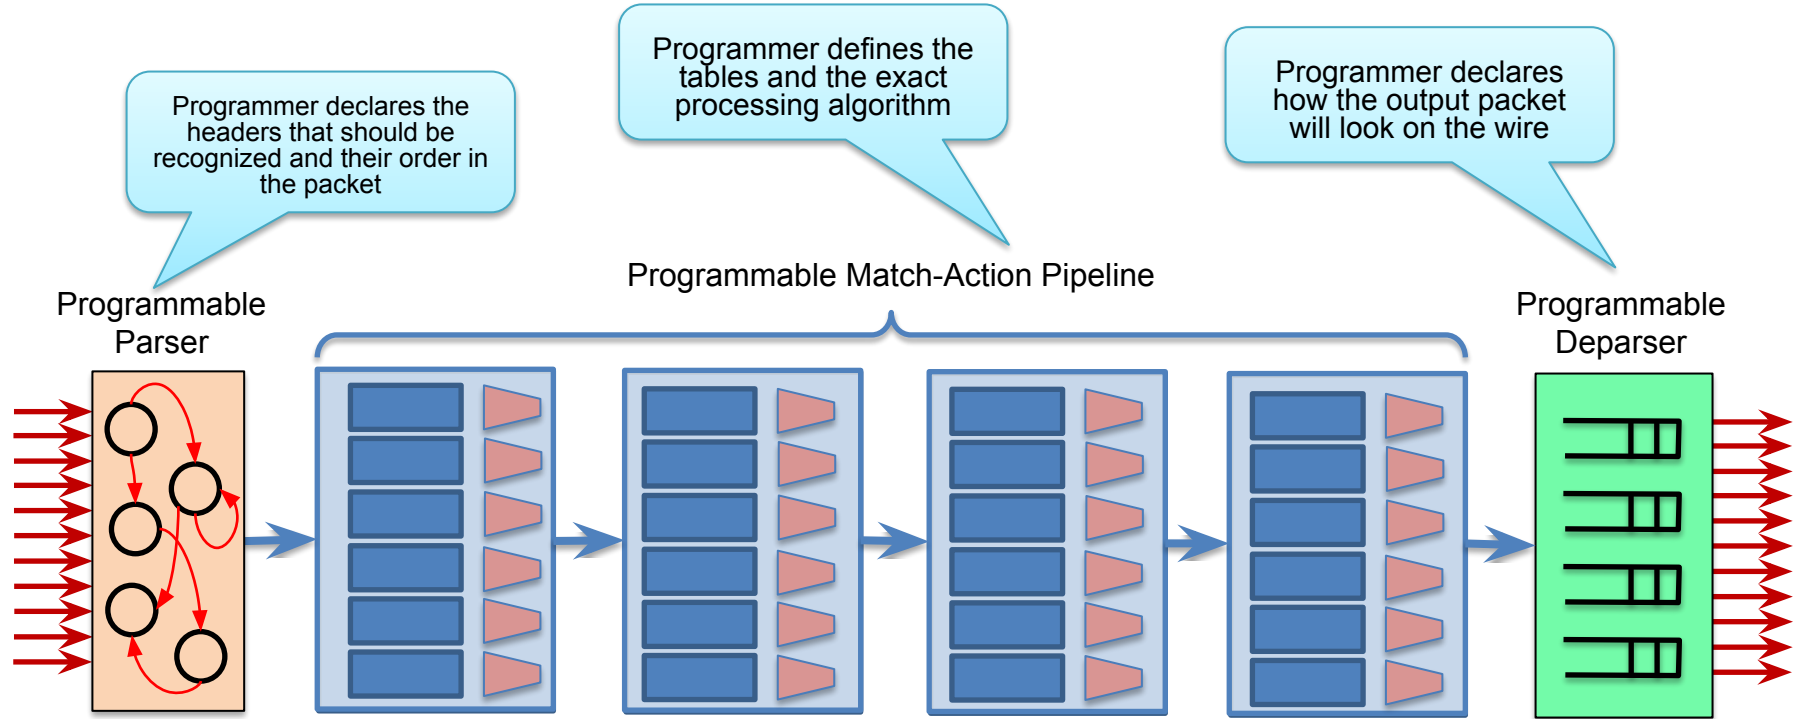
- **Clearly defined semantics**
 - You can describe what your data plane program is doing
- **Expressive**
 - Supports a wide range of architectures through standard methodology
- **High-level, Target-independent**
 - Uses conventional constructs
 - Compiler manages the resources and deals with the hardware
- **Type-safe**
 - Enforces good software design practices and eliminates “stupid” bugs
- **Agility**
 - High-speed networking devices become as flexible as any software
- **Insight**
 - Freely mixing packet headers and intermediate results



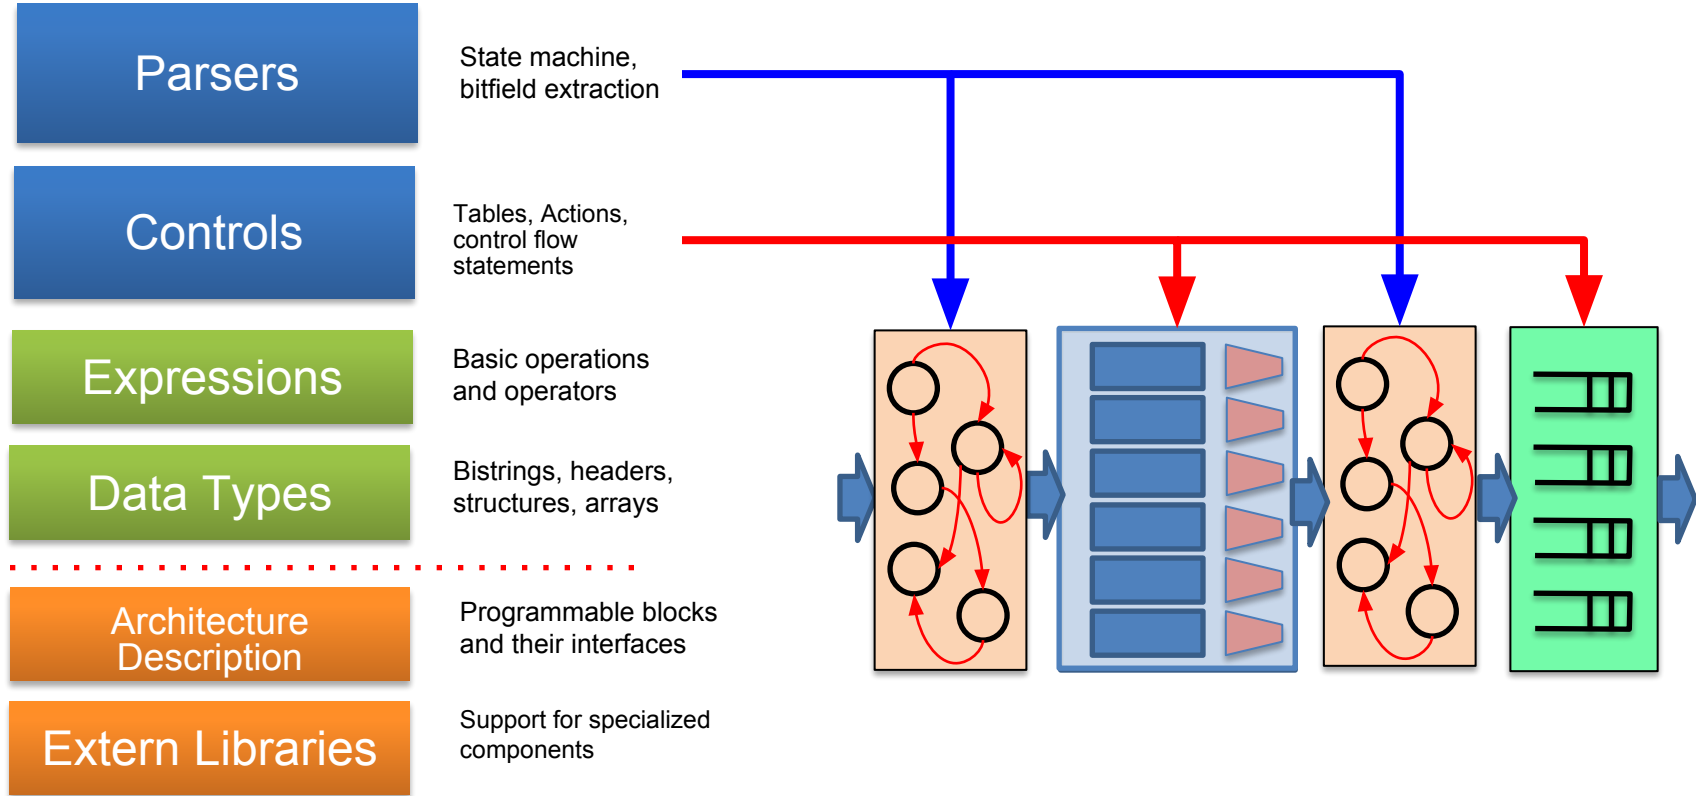
P4_16 Data Plane Model



PISA: Protocol-Independent Switch Architecture



P4₁₆ Language Elements



P4_16 Approach

Term	Explanation
P4 Target	An embodiment of a specific hardware implementation
P4 Architecture	Provides an interface to program a target via some set of P4-programmable components, externs, fixed components

