


Disclaimer

- These are not slides from DTU (refer to p4.org)
- You can find the original slides at:
 - https://github.com/p4lang/tutorials/blob/master/P4_tutorial.pdf
- We removed some of the slides and organized in several groups
- You are welcome to check the original slides for a wider perspective
- If you have doubts about:
 - What P4 is, why it is beneficial or the core concepts of P4: Refer to file “**1 - P4 Information.pdf**”
 - What the core elements of P4 programming are (controls, externs, tables, data types), how to program a target, how the architecture refers to a P4 program and then: Refer to file “**2 - Basics on P4 programming.pdf**”
 - What P4 runtime is, how it related to P4 overall and what you can do with it: Refer to file “**3 - P4 Runtime.pdf**”



P4 Language Tutorial

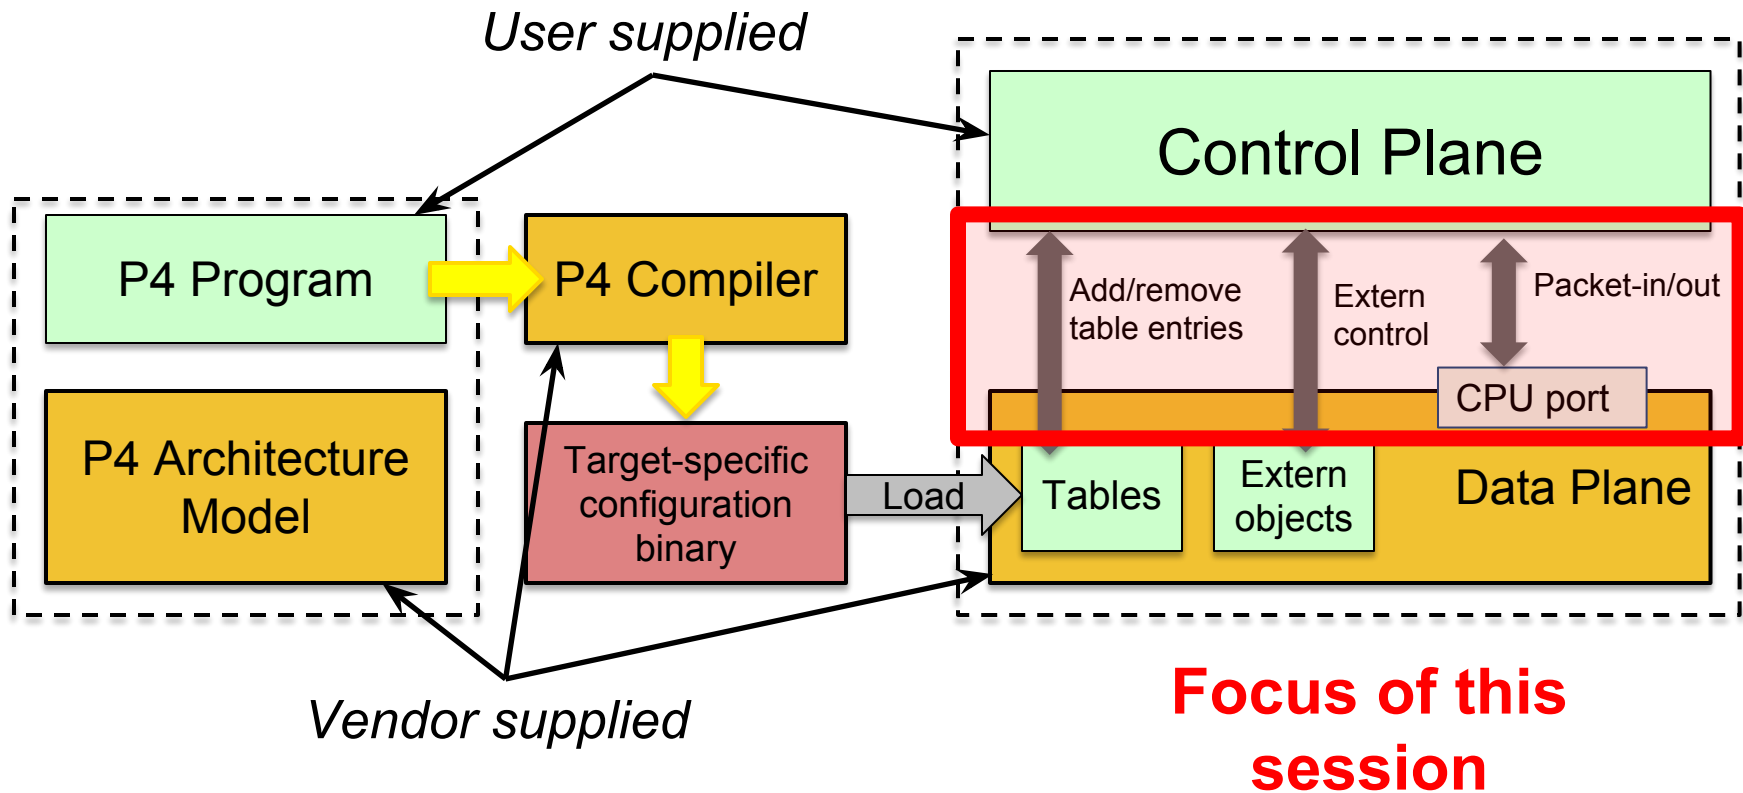


P4Runtime

- **API overview**
- **Workflow**
- **Exercise - Tunneling**



Runtime control of P4 data planes

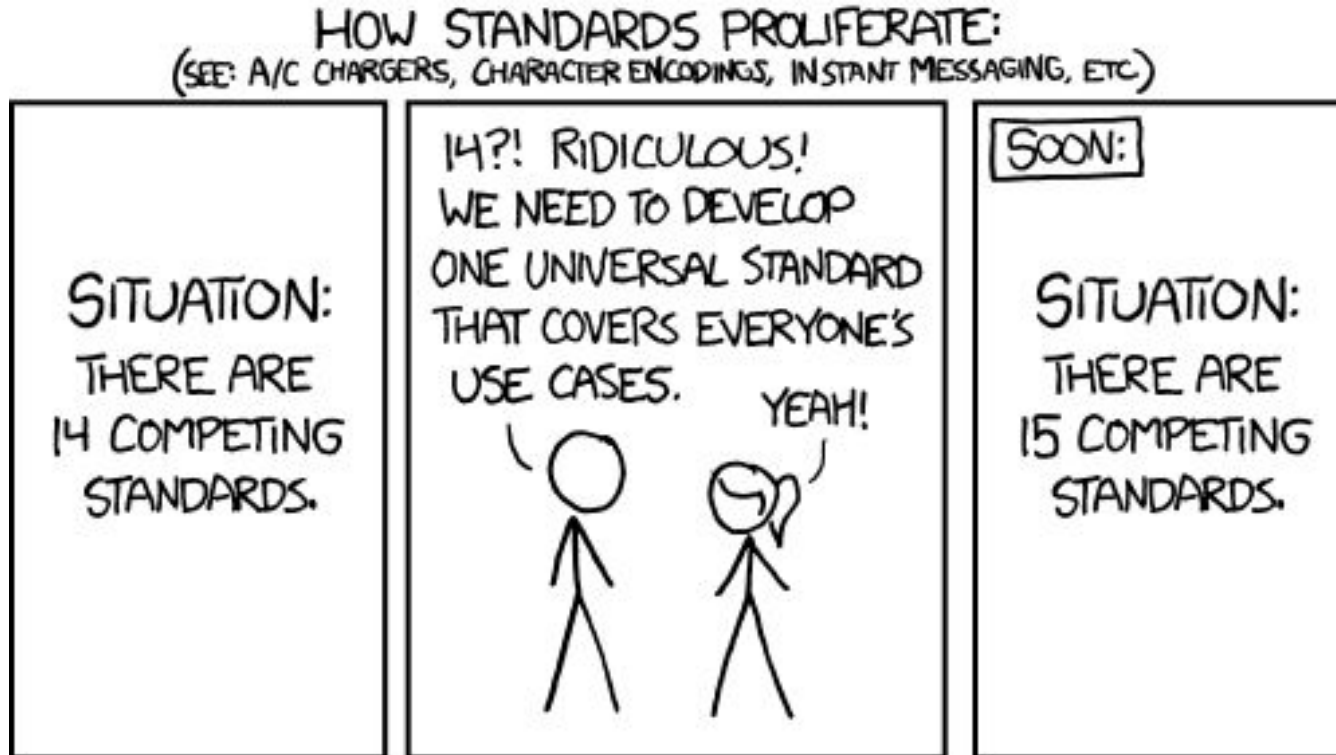


Existing approaches to runtime control

- **P4 compiler auto-generated runtime APIs**
 - Program-dependent -- hard to provision new P4 program without restarting the control plane!
- **BMv2 CLI**
 - Program-independent, but target-specific -- control plane not portable!
- **OpenFlow**
 - Target-independent, but protocol-dependent -- protocol headers and actions baked in the specification!
- **OCP Switch Abstraction Interface (SAI)**
 - Target-independent, but protocol-dependent



Why do we need another data plane control API?

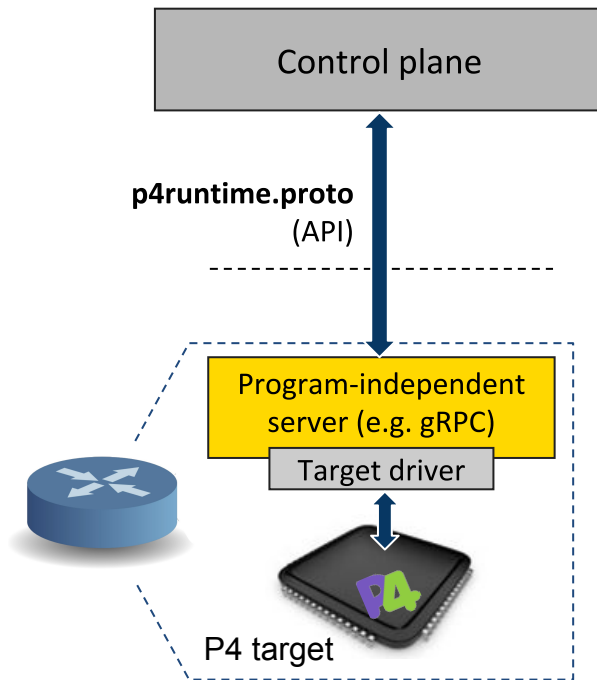


Properties of a runtime control API

| API | Target-independent | Protocol-independent |
|----------------------------|--------------------|----------------------|
| P4 compiler auto-generated | ✓ | ✗ |
| BMv2 CLI | ✗ | ✓ |
| OpenFlow | ✓ | ✗ |
| SAI | ✓ | ✗ |
| P4Runtime | ✓ | ✓ |

What is P4Runtime?

- **Framework for runtime control of P4 targets**
 - Open-source API + server implementation
<https://github.com/p4lang/PI>
 - Initial contribution by Google and Barefoot
- **Work-in-progress by the p4.org API WG**
 - Draft of version 1.0 available
- **Protobuf-based API definition**
 - p4runtime.proto
 - gRPC transport
- **P4 program-independent**
 - API doesn't change with the P4 program
- **Enables field-reconfigurability**
 - Ability to push new P4 program without recompiling the software stack of target switches



Protocol Buffers (protobuf) Basics

- Language for describing data for serialization in a structured way
- Common binary wire-format
- Language-neutral
 - Code generators for: *Action Script, C, C++, C#, Clojure, Lisp, D, Dart, Erlang, Go, Haskell, Java, Javascript, Lua, Objective C, OCaml, Perl, PHP, Python, Ruby, Rust, Scala, Swift, Visual Basic, ...*
- Platform-neutral
- Extensible and backwards compatible
- Strongly typed

```
syntax = "proto3";

message Person {
  string name = 1;
  int32 id = 2;
  string email = 3;

  enum PhoneType {
    MOBILE = 0;
    HOME = 1;
    WORK = 2;
  }


  message PhoneNumber {
    string number = 1;
    PhoneType type = 2;
  }

  repeated PhoneNumber phone = 4;
}
```



gRPC Basics

- Use Protocol Buffers to define service API and messages
- Automatically generate native stubs in:
 - C / C++
 - C#
 - Dart
 - Go
 - Java
 - Node.js
 - PHP
 - Python
 - Ruby
- Transport over HTTP/2.0 and TLS
 - Efficient single TCP connection implementation that supports bidirectional streaming



gRPC Service Example

```
// The greeter service definition.
service Greeter {
    // Sends a greeting
    rpc SayHello (HelloRequest) returns (HelloReply) {}
}

// The request message containing the user's name.
message HelloRequest {
    string name = 1;
}

// The response message containing the greetings
message HelloReply {
    string message = 1;
}
```

More details here: <https://grpc.io/docs/guides/>



P4Runtime Service

Enables a local or remote entity to arbitrate mastership, load the pipeline/program, send/receive packets, and read and write forwarding table entries, counters, and other P4 entities.

```
service P4Runtime {  
  rpc Write(WriteRequest) returns (WriteResponse) {}  
  rpc Read(ReadRequest) returns (stream ReadResponse) {}  
  rpc SetForwardingPipelineConfig(SetForwardingPipelineConfigRequest)  
    returns (SetForwardingPipelineConfigResponse) {}  
  rpc GetForwardingPipelineConfig(GetForwardingPipelineConfigRequest)  
    returns (GetForwardingPipelineConfigResponse) {}  
  rpc StreamChannel(stream StreamMessageRequest)  
    returns (stream StreamMessageResponse) {}  
}
```



P4Runtime Service

P4Runtime Protobuf Definition:

<https://github.com/p4lang/p4runtime/blob/master/proto/p4/v1/p4runtime.proto>

Service Specification:

Working draft of version 1.0 is available now

<https://p4.org/p4-spec/docs/P4Runtime-v1.0.0.pdf>



P4Runtime Write Request

```
message WriteRequest {  
  uint64 device_id = 1;  
  uint64 role_id = 2;  
  Uint128 election_id = 3;  
  repeated Update updates = 4;  
}
```

```
message Update {  
  enum Type {  
    UNSPECIFIED = 0;  
    INSERT = 1;  
    MODIFY = 2;  
    DELETE = 3;  
  }  
  Type type = 1;  
  Entity entity = 2;  
}
```

```
message Entity {  
  oneof entity {  
    ExternEntry extern_entry = 1;  
    TableEntry table_entry = 2;  
    ActionProfileMember  
      action_profile_member = 3;  
    ActionProfileGroup  
      action_profile_group = 4;  
    MeterEntry meter_entry = 5;  
    DirectMeterEntry direct_meter_entry = 6;  
    CounterEntry counter_entry = 7;  
    DirectCounterEntry direct_counter_entry = 8;  
    PacketReplicationEngineEntry  
      packet_replication_engine_entry = 9;  
    ValueSetEntry value_set_entry = 10;  
    RegisterEntry register_entry = 11;  
  }  
}
```



P4Runtime Table Entry

p4runtime.proto simplified excerpts:

```
message TableEntry {  
  uint32 table_id;  
  repeated FieldMatch match;  
  Action action;  
  int32 priority;  
  ...  
}
```

```
message Action {  
  uint32 action_id;  
  message Param {  
    uint32 param_id;  
    bytes value;  
  }  
  repeated Param params;  
}
```

```
message FieldMatch {  
  uint32 field_id;  
  message Exact {  
    bytes value;  
  }  
  message Ternary {  
    bytes value;  
    bytes mask;  
  }  
  ...  
  oneof field_match_type {  
    Exact exact;  
    Ternary ternary;  
    ...  
  }  
}
```

To add a table entry, the control plane needs to know:

- **IDs of P4 entities**
 - Tables, field matches, actions, params, etc.
- **Field matches for the particular table**
 - Match type, bitwidth, etc.
- **Parameters for the particular action**
- **Other P4 program attributes**

Full protobuf definition:

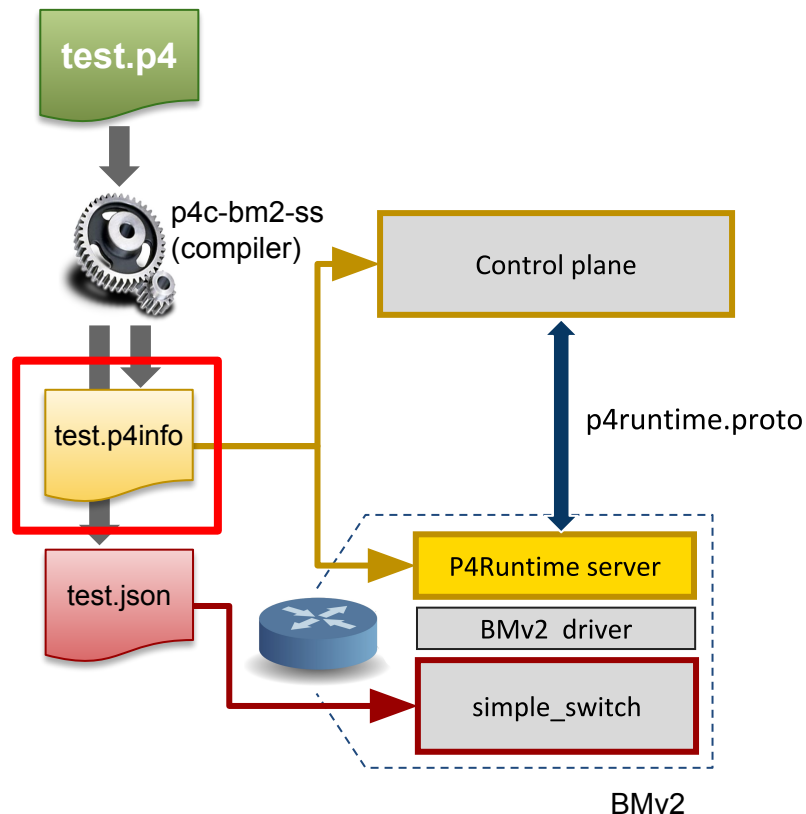
<https://github.com/p4lang/PI/blob/master/proto/p4/p4runtime.proto>



P4Runtime workflow

P4Info

- **Captures P4 program attributes needed at runtime**
 - IDs for tables, actions, params, etc.
 - Table structure, action parameters, etc.
- **Protobuf-based format**
- **Target-independent compiler output**
 - Same P4Info for BMv2, ASIC, etc.



Full P4Info protobuf specification:

<https://github.com/p4lang/p4runtime/blob/master/proto/p4/config/v1/p4info.proto>



P4Info example

basic_router.p4

```
...  
  
action ipv4_forward(bit<48> dstAddr,  
                    bit<9> port) {  
    /* Action implementation */  
}  
  
...  
  
table ipv4_lpm {  
    key = {  
        hdr.ipv4.dstAddr: lpm;  
    }  
    actions = {  
        ipv4_forward;  
        ...  
    }  
    ...  
}
```



P4 compiler

basic_router.p4info

```
actions {  
    id: 16786453  
    name: "ipv4_forward"  
    params {  
        id: 1  
        name: "dstAddr"  
        bitwidth: 48  
        ...  
        id: 2  
        name: "port"  
        bitwidth: 9  
    }  
}  
...  
tables {  
    id: 33581985  
    name: "ipv4_lpm"  
    match_fields {  
        id: 1  
        name: "hdr.ipv4.dstAddr"  
        bitwidth: 32  
        match_type: LPM  
    }  
    action_ref_id: 16786453  
}
```

P4Runtime Table Entry Example

basic_router.p4


```
action ipv4_forward(bit<48> dstAddr,  
                    bit<9> port) {  
    /* Action implementation */  
}  
  
table ipv4_lpm {  
    key = {  
        hdr.ipv4.dstAddr: lpm;  
    }  
    actions = {  
        ipv4_forward;  
        ...  
    }  
    ...  
}
```



Logical view of table entry

```
hdr.ipv4.dstAddr=10.0.1.1/32  
-> ipv4_forward(00:00:00:00:00:10, 7)
```

Control plane generates



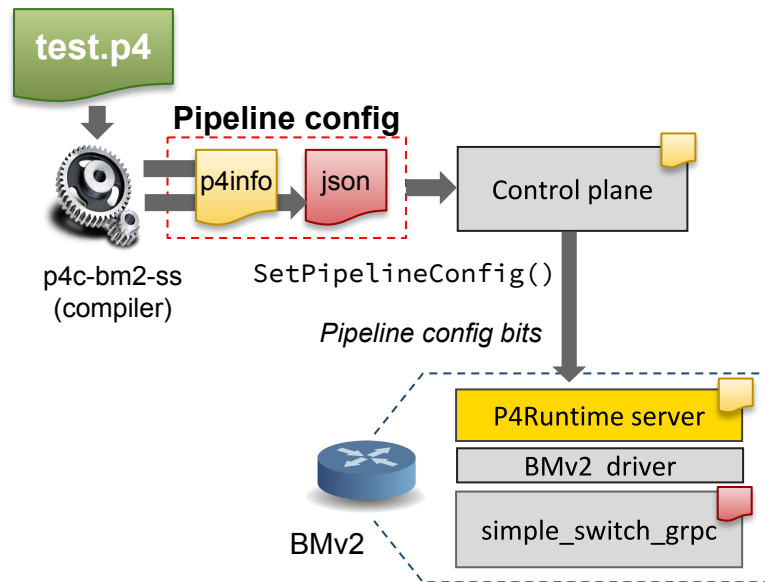
Protobuf message

```
table_entry {  
  table_id: 33581985  
  match {  
    field_id: 1  
    lpm {  
      value: "\n\000\001\001"  
      prefix_len: 32  
    }  
  }  
  action {  
    action_id: 16786453  
    params {  
      param_id: 1  
      value: "\000\000\000\000\000\000\n"  
    }  
    params {  
      param_id: 2  
      value: "\000\007"  
    }  
  }  
}
```



P4Runtime SetPipelineConfig

```
message SetForwardingPipelineConfigRequest {
  enum Action {
    UNSPECIFIED = 0;
    VERIFY = 1;
    VERIFY_AND_SAVE = 2;
    VERIFY_AND_COMMIT = 3;
    COMMIT = 4;
    RECONCILE_AND_COMMIT = 5;
  }
  uint64 device_id = 1;
  uint64 role_id = 2;
  Uint128 election_id = 3;
  Action action = 4;
  ForwardingPipelineConfig config = 5;
}
```



```
message ForwardingPipelineConfig {
  config.P4Info p4info = 1;
  // Target-specific P4 configuration.
  bytes p4_device_config = 2;
}
```



P4Runtime StreamChannel

```
message StreamMessageRequest {  
  oneof update {  
    MasterArbitrationUpdate  
      arbitration = 1;  
    PacketOut packet = 2;  
    DigestListAck digest_ack = 3;  
  }  
}
```

```
message StreamMessageResponse {  
  oneof update {  
    MasterArbitrationUpdate  
      arbitration = 1;  
    PacketIn packet = 2;  
    DigestList digest = 3;  
  }  
}
```

// Packet sent from the controller to the switch.

```
message PacketOut {  
  bytes payload = 1;  
  // This will be based on P4 header annotated as  
  // @controller_header("packet_out").  
  // At most one P4 header can have this annotation.  
  repeated PacketMetadata metadata = 2;  
}
```

// Packet sent from the switch to the controller.

```
message PacketIn {  
  bytes payload = 1;  
  // This will be based on P4 header annotated as  
  // @controller_header("packet_in").  
  // At most one P4 header can have this annotation.  
  repeated PacketMetadata metadata = 2;  
}
```



P4Runtime Common Parameters

- **device_id**
 - Specifies the specific forwarding chip or software bridge
 - **Set to 0 for single chip platforms**
- **role_id**
 - Corresponds to a role with specific capabilities (i.e. what operations, P4 entities, behaviors, etc. are in the scope of a given role)
 - Role definition is currently agreed upon between control and data planes offline
 - **Default role_id (0) has full pipeline access**
- **election_id**
 - P4Runtime supports mastership on a per-role basis
 - Client with the highest election ID is referred to as the "master", while all other clients are referred to as "slaves"
 - **Set to 0 for single instance controllers**

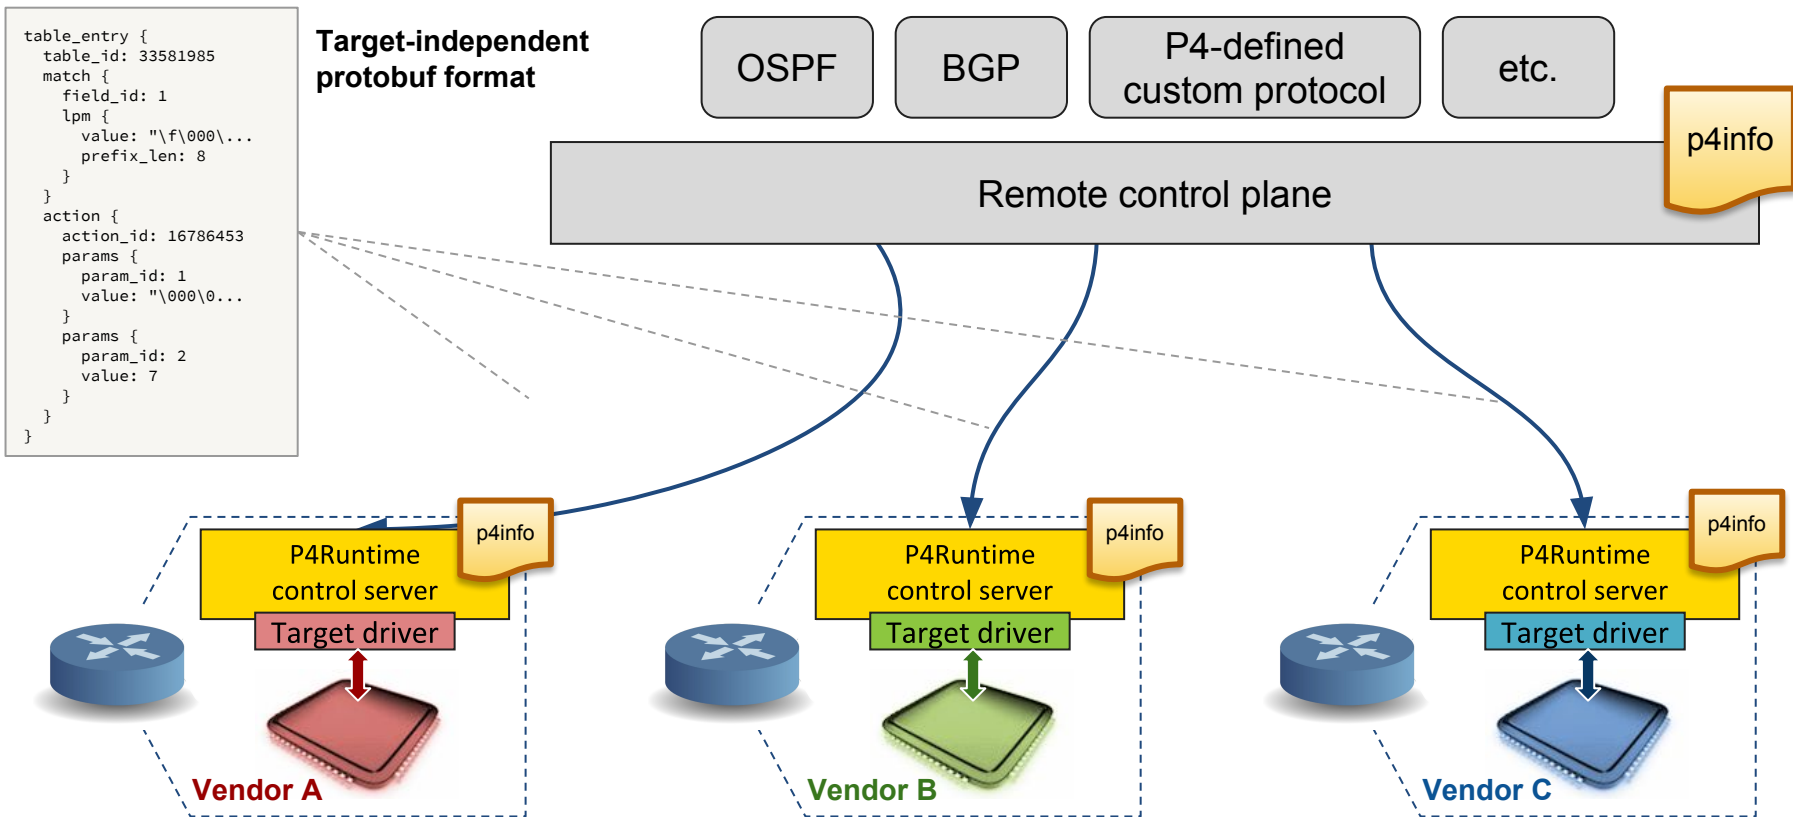


Mastership Arbitration

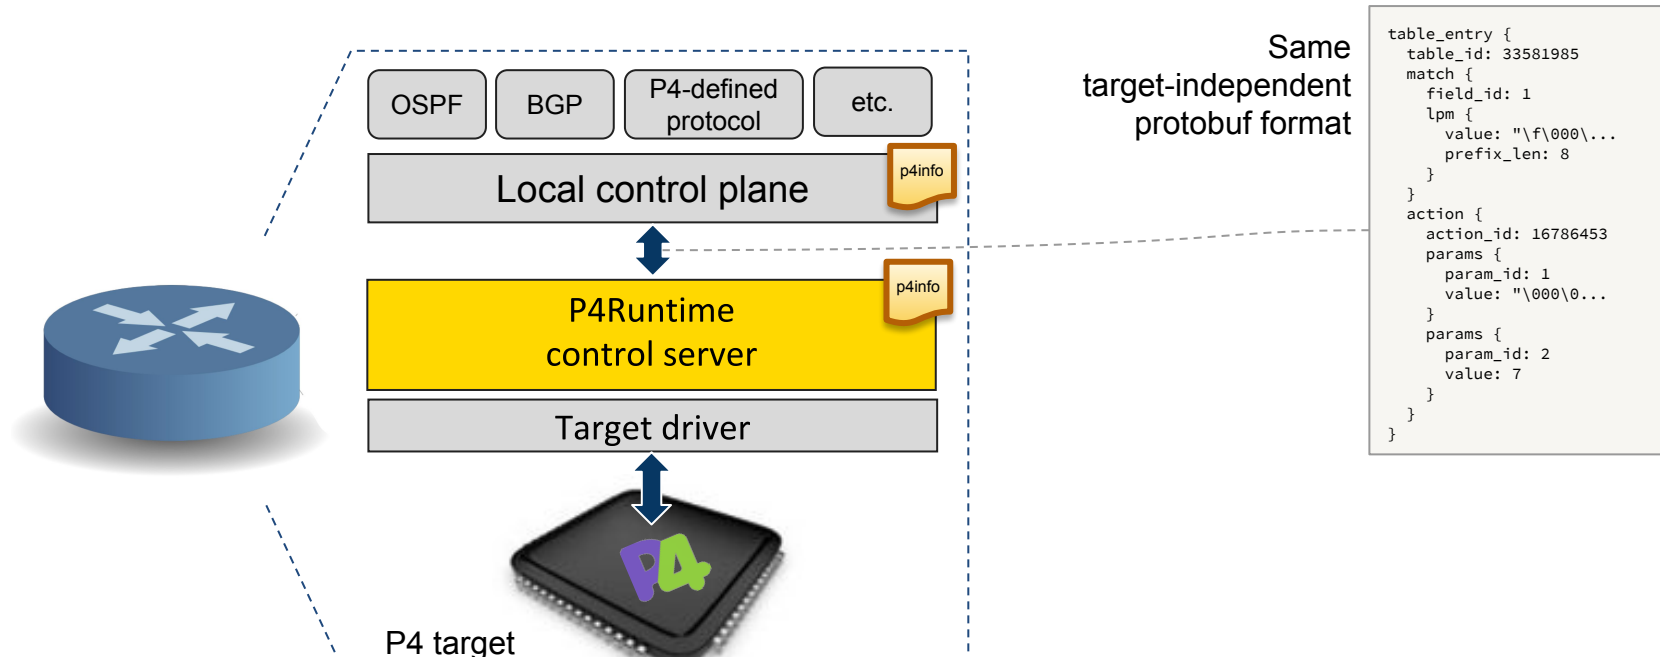
- **Upon connecting to the device, the client (e.g. controller) needs to open a StreamChannel**
- **The client must advertise its `role_id` and `election_id` using a `MasterArbitrationUpdate` message**
 - If `role_id` is not set, it implies the default role and will be granted full pipeline access
 - The `election_id` is opaque to the server and determined by the control plane (can be omitted for single-instance control plane)
- **The switch marks the client for each role with the highest `election_id` as master**
- **Master can:**
 - Perform Write requests
 - Receive PacketIn messages
 - Send PacketOut messages



Remote control



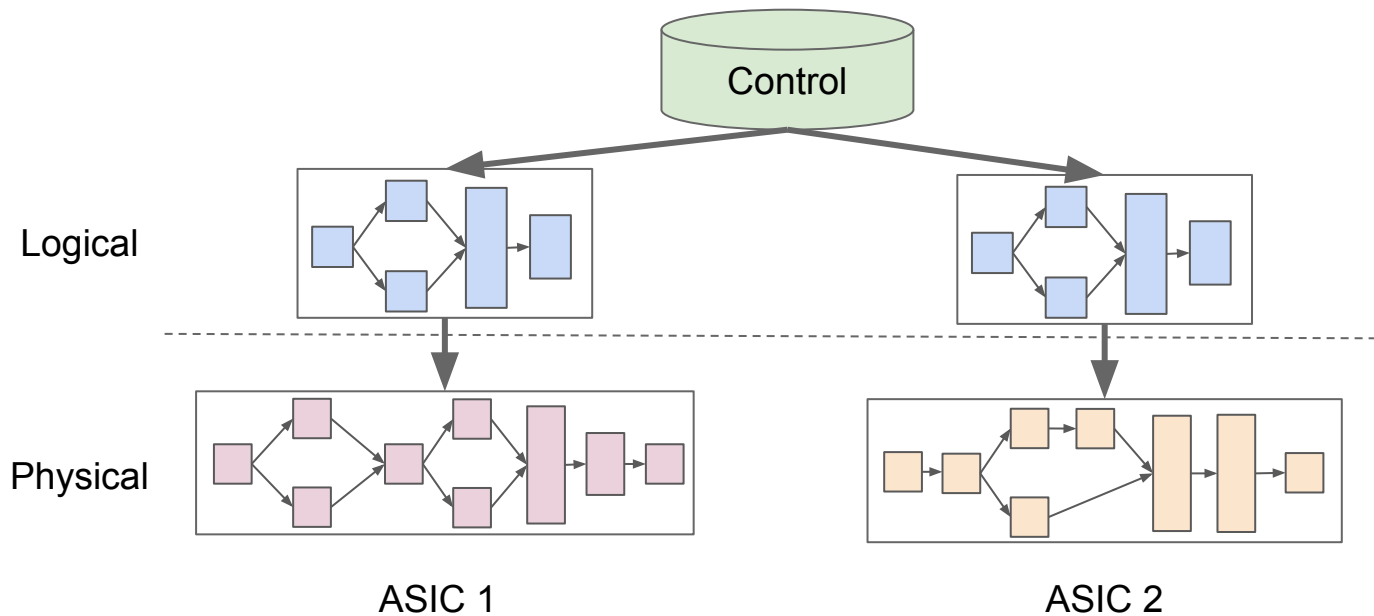
Local control



The P4Runtime API can be used equally well
by a remote or local control plane

P4 Program as Fixed-Function Chip Abstraction

- P4 program tailored to apps / role - does not describe the hardware
- Switch maps program to fixed-function ASIC
- Enables portability



P4Runtime API recap

Things we covered:

- **P4Info**
- **Table entries**
- **Set pipeline config**

What we didn't cover:

- **How to control other P4 entities**
 - Externs, counters, meters
- **Packet-in/out support**
- **Controller replication**
 - Via master-slave arbitration
- **Batched reads/writes**
- **Switch configuration**
 - Outside the P4Runtime scope
 - Achieved with other mechanisms
 - e.g., OpenConfig and gNMI

Work-in-progress by the p4.org API WG
Expect API changes in the future



P4Runtime workflow



P4 compiler generates 2 files:

1. Target-specific binaries
 - Used to configure switch pipeline (e.g. binary config for ASIC, bitstream for FPGA, etc.)
2. P4Info file
 - Captures P4 program attributes needed to runtime control
 - Tables, actions, parameters, etc.
 - Protobuf-based format
 - Target-independent compiler output
 - Same P4Info for SW switch, ASIC, etc.

Full P4Info protobuf specification:

<https://github.com/p4lang/PI/blob/master/proto/p4/config/p4info.proto>

