


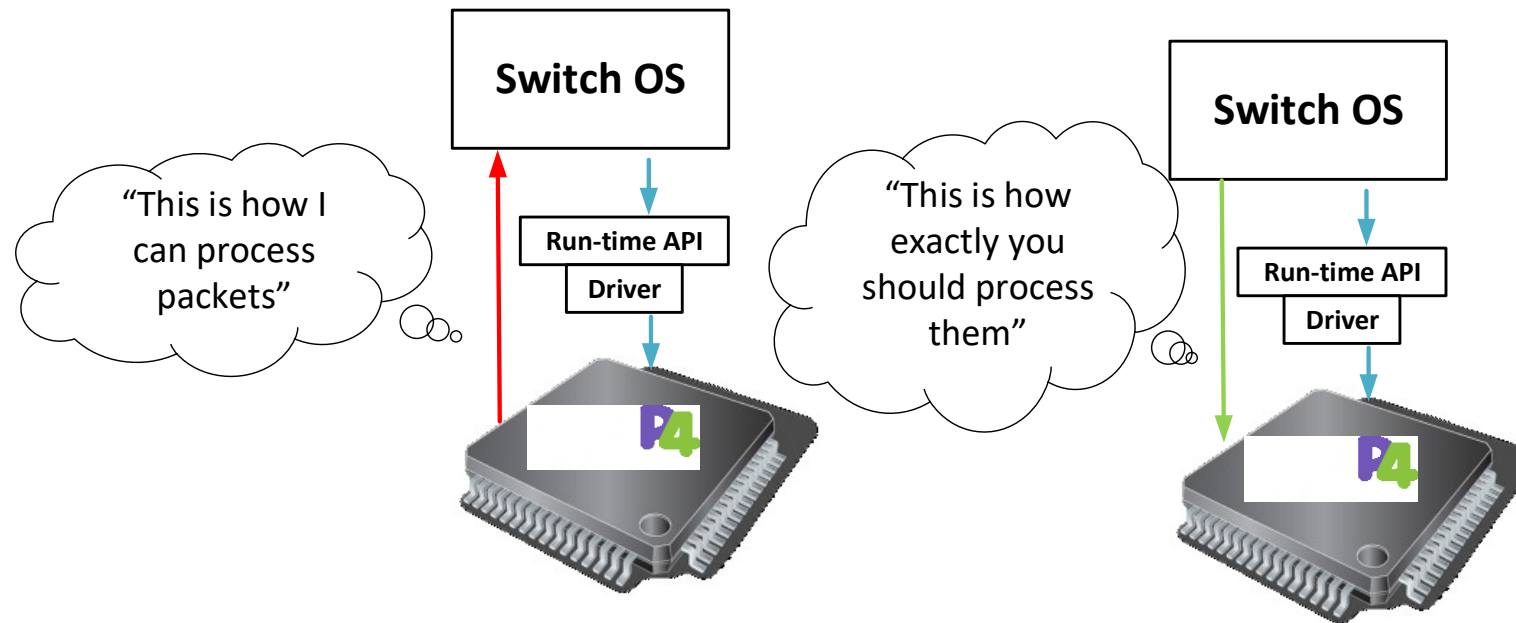
# A short intro to P4

Eder Ollora Zaballa



# Motivation


- SDN related: Each Openflow specification added more and more headers. Instead of incrementing the support for new headers in each specification, why not let programmers define the data plane?



# Terminology

- This is a very important part since you will find a lot of information, typically mixing terms:
- **Target:** The embodiment of a specific hardware implementation.
  - For instance: **Simple Switch**. So Simple Switch is a target and is built on top of a set of tools named Behavioral Model v2 (aka Bmv2) [1]. Sometimes people refer to their switch as a Bmv2 switch but in reality you can implement another target based on Bmv2 [2].
- **Architecture:** The interface to program a target, and defines a set of P4-programmable components, externs or even fixed components.
  - For instance: **V1 model** [3]. The architecture defines which blocks we can program.

# V1 Model (Architecture)



# V1 Model (Architecture)

```

/*****
*****  P A R S E R  *****/
*****/

```

```

parser MyParser(packet_in packet,
out headers hdr,
inout metadata meta,
inout standard_metadata_t standard_metadata) {

```

```

state start {
/* TODO: add parser logic */
transition accept;
}
}

```

```

/*****
*****  C H E C K S U M   V E R I F I C A T I O N  *****/
*****/

```

```

control MyVerifyChecksum(inout headers hdr, inout metadata meta) {
apply { }
}

```

```

/*****
*****  I N G R E S S   P R O C E S S I N G  *****/
*****/

```

```

control MyIngress(inout headers hdr,
inout metadata meta,
inout standard_metadata_t standard_metadata) {
action drop() {
mark_to_drop();
}

```

```

action ipv4_forward(macaddr_t dstaddr, egressSpec_t port) {
/* TODO: fill out code in action body */
}

```

```

table ipv4_lpm {
key = {
hdr.ipv4.dstaddr: lpm;
}
actions = {
ipv4_forward;
drop;
NoAction;
}
size = 1024;
default_action = NoAction();
}

```

```

apply {
/* TODO: fix ingress control logic
* - ipv4_lpm should be applied only when IPv4 header is valid
*/
ipv4_lpm.apply();
}
}

```

```

parser Parser<H, M>(packet_in b,
out H parsedHdr,
inout M meta,
inout standard_metadata_t standard_metadata);

```

```

/* The only legal statements in the implementation of the
VerifyChecksum control are: block statements, calls to the
verify_checksum and verify_checksum_with_payload methods,
and return statements. */

```

```

control VerifyChecksum<H, M>(inout H hdr,
inout M meta);

```

```

@pipeline
control Ingress<H, M>(inout H hdr,
inout M meta,
inout standard_metadata_t standard_metadata);

```

```

@pipeline
control Egress<H, M>(inout H hdr,
inout M meta,
inout standard_metadata_t standard_metadata);

```

```

/* The only legal statements in the implementation of the
ComputeChecksum control are: block statements, calls to the
update_checksum and update_checksum_with_payload methods,
and return statements. */
control ComputeChecksum<H, M>(inout H hdr,
inout M meta);

```

```

@deparser
control Deparser<H>(packet_out b, in H hdr);

```

```

package V1Switch<H, M>(Parser<H, M> p,
VerifyChecksum<H, M> vr,
Ingress<H, M> ig,
Egress<H, M> eg,
ComputeChecksum<H, M> ck,
Deparser<H> dep
);

```

```

/*****
*****  E G R E S S   P R O C E S S I N G  *****/
*****/

```

```

control MyEgress(inout headers hdr,
inout metadata meta,
inout standard_metadata_t standard_metadata) {
apply { }
}

```

```

/*****
*****  C H E C K S U M   C O M P U T A T I O N  *****/
*****/

```

```

control MyComputeChecksum(inout headers hdr, inout metadata meta) {
apply {
update_checksum(
hdr.ipv4.isValid(),
{ hdr.ipv4.version,
hdr.ipv4.ihl,
hdr.ipv4.diffserv,
hdr.ipv4.totallen,
hdr.ipv4.identification,
hdr.ipv4.flags,
hdr.ipv4.fragoffset,
hdr.ipv4.ttl,
hdr.ipv4.protocol,
hdr.ipv4.srcaddr,
hdr.ipv4.dstaddr },
hdr.ipv4.hdrChecksum,
HashAlgorithm.csum16);
}
}

```

```

/*****
*****  D E P A R S E R  *****/
*****/

```

```

control MyDeparser(packet_out packet, in headers hdr) {
apply {
/* TODO: add deparser logic */
}
}

```

```

/*****
*****  S W I T C H  *****/
*****/

```

```

V1Switch(
MyParser(),
MyVerifyChecksum(),
MyIngress(),
MyEgress(),
MyComputeChecksum(),
MyDeparser()
) main;

```

# Standard metadata (V1 model)

[4]

```
struct standard_metadata_t {
    bit<9> ingress_port;
    bit<9> egress_spec;
    bit<9> egress_port;
    bit<32> clone_spec;
    bit<32> instance_type;
    bit<1> drop;
    bit<16> recirculate_port;
    bit<32> packet_length;
    bit<32> enq_timestamp;
    bit<19> enq_qdepth;
    bit<32> deq_timedelta;
    bit<19> deq_qdepth;
    bit<48> ingress_global_timestamp;
    bit<32> lf_field_list;
    bit<16> mcast_grp;
    bit<1> resubmit_flag;
    bit<16> egress_rid;
    bit<1> checksum_error;
}
```

**ingress\_port** - the port on which the packet arrived

**egress\_spec** - the port to which the packet should be sent to

**egress\_port** - the port on which the packet is departing from (read only in egress pipeline)

# Target programming

